

Christian Mitton

I worked by myself.

Summary

My code was built off of the foundation provided in `tfs_mkfs()`, since that is where the framework of my file system is created. So in order to illustrate my thought process I'll discuss `tfs_mkfs()` first and how it creates the file system, and then I'll explain how I envisioned the provided functions would interact with my file system.

`tfs_mkfs()`:

Step 1.) Create `super block (sb)`, then write it to the disk

- The super block is created and all its fields are filled out. In the `super block`, the number that corresponds to the `i_bitmap_blk`, `d_bitmap_blk`, `i_start_blk` and `d_start_blk` refers to the index they will be located in on the disk. After the values are assigned, the `superblock` is written to the 0th index in the disk file using `bio_write()`.

Step 2.) Create bitmaps, set 0th index in both bitmaps for root, write bitmaps to disk

- The inode and data bitmaps are created, and then their 0th index is set to 1 for the root. After this is done, they are both written to the disk.

Step 3.) Create Inode and dirent for root, write both to disk

- The inode and dirent for the root is created, the fields are filled out and its ino is set to 0. The inode for the root is written to the disk at its corresponding index (`i_start_blk`) and the dirent for root is written at its corresponding index (`d_start_blk`)

Step 4.) Fill inode region with available inodes

- I considered the inode region to start at the `i_start_blk` index and end at the `d_start_blk - 1` index. The data region is everything after. This region of the disk is filled with free inodes. The number of inodes possible was calculated using this formula:

```
spaceNeededForInodes = (sizeof(struct inode) * MAX_INUM)
numOfInodes = spaceNeededForInodes / BLOCK_SIZE
```

When the above steps are completed and `tfs_mkfs()` finished, the file "disk" should have the following structure:

DISK						
			Inode Region		Data Region	
0	1	2	3	...	69	...
super block	i_bitmap_blk	d_bitmap_blk	i_start_blk	...	d_start_blk	...
Information about disk	100000000..	1000000000..	root inode	remaining possible inodes	root dirent	remaining possible datablocks

get_avail_ino():

This function traverses the `i_bitmap_blk` to find the next available index (indicated by a 0). Once it finds this index, it sets that index to 1 in the `i_bitmap_blk`, writes this updated bitmap to the disk, and then it returns the block number of the next available inode in the disk.

get_avail_blkno():

This function does the exact same as `get_avail_ino()`, only it traverses the `d_bitmap_blk` instead, and finds the next available data block on the disk.

`readi():`

Given an ino number and a pointer to an inode, this function will get the inode on the disk that corresponds to the ino, and then it'll load that inode into the specified pointer.

`writei():`

This function is similar to `readi()`, only instead of loading an inode from the disk into the specified pointer, it'll overwrite the inode on the disk with the specified inode.

`dir_find():`

Given a name (`f_name`), this function will perform all the necessary steps to extract the dirent corresponding to the name on the disk, and load it to the provided dirent pointer.

`int dir_add():`

Given a name (`f_name`), inode number (`ino`) and a current directory (`dir_inode`) as arguments, this function will add a new dirent entry to the current directory's `direct_ptr` array. This entry will have the inode number that was passed into the function. It'll then write this new dirent to the data block region on the disk, and the name of this dirent will be assigned to `f_name`.

`Int dir_remove():`

This function will search the current directory (`dir_inode`), and if this directory contains a dirent with the name "`fname`" it will remove it by setting it's value to 0 in the current directory's `direct_ptr` array. It will also set the validity of the inode for the dirent to remove.

int get_node_by_path():

This is a recursive function that given a path, say: `"/folder1/folder2/file.txt"`, it will load the inode of `file.txt` into "`inode`". This function assumes that the path will always start from root. There are many steps involved, but to summerize briefly:

1.) The path is split into tokens:

- `"/folder1/folder2/file.txt" ==> "." "folder1" "folder2" "file.txt"`

2.) Each token is passed into the recursive function `getPathRec()` as "`currentDirName`."

3.) With each recursive call, if `currentDirName` is a directory, it searches for the next token in the path

4.) If the last token is found in it's parent directory it returns an inode to `get_node_by_path()`, and this inode is read into the original "`inode`" argument that was passed to `get_node_by_path()`.

5.) If the last token is not found, the file doesn't exist and `getPathRec()` returns NULL. (The rest of the code assumes the path is always valid)